

```
In [21]: import numpy as np
import pandas as pd
# from mnist import MNIST
from sklearn.metrics import confusion_matrix
from sklearn.mixture import GaussianMixture
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
%matplotlib inline

In [22]: # #Generating Data Files
# mndata = MNIST('')

In [23]: # imgs, labels = mndata.load_training()

In [24]: # train_data = pd.DataFrame(imgs)
# train_labels = pd.DataFrame(list(labels))

In [25]: # imgs, labels = mndata.load_testing()
# test_data = pd.DataFrame(imgs)
# test_labels = pd.DataFrame(list(labels))

In [26]: # train_data_subset = train_data.loc[:5999]
# train_labels_subset = train_labels.loc[:5999]
# test_data_subset = test_data.loc[:999]
# test_labels_subset = test_labels.loc[:999]

In [27]: # train_data_subset = train_data_subset/255.0
# test_data_subset = test_data_subset/255.0
# train_data_subset = train_data_subset.round()
# test_data_subset = test_data_subset.round()

In [28]: # train_data_subset.to_csv('mnist_train_data.csv', index=False, header=False)
# train_labels_subset.to_csv('mnist_train_labels.csv', index=False, header=False)
# test_data_subset.to_csv('mnist_test_data.csv', index=False, header=False)
# test_labels_subset.to_csv('mnist_test_labels.csv', index=False, header=False)

In [54]: train_data = pd.read_csv('mnist_train_data.csv', header=None)
train_labels = pd.read_csv('mnist_train_labels.csv', names=['label'])
test_data = pd.read_csv('mnist_test_data.csv', header=None)
test_labels = pd.read_csv('mnist_test_labels.csv', names=['label'])

In [34]: #Question 1
# 1 mark: correct answer
prior = train_labels.groupby(['label'])['label'].count().reset_index(drop=True)/6000
```

```
In [35]: prior
```

```
Out[35]: 0    0.098667
         1    0.111833
         2    0.096833
         3    0.101333
         4    0.103833
         5    0.085667
         6    0.101333
         7    0.108500
         8    0.091833
         9    0.100167
         Name: label, dtype: float64
```

```
In [36]: #Question 2
         # -1 mark: no laplace smoothing
         # -1 mark: wrong answer, correct approach
         # -3 marks: wrong approach, no answer

         data = pd.concat([train_data, train_labels], axis=1)
         num_data = data.groupby(['label']).sum().reset_index()
         del num_data['label']
         num_data = num_data+1
```

```
In [37]: denom_data = data.groupby(['label'])['label'].count().reset_index(drop=True)
         denom_data = denom_data + 2
         likelihood = num_data.divide(denom_data, axis=0)
```

```
In [38]: np.log(likelihood.max(axis=1))
```

```
Out[38]: 0    -0.160343
         1    -0.014970
         2    -0.316098
         3    -0.212950
         4    -0.162990
         5    -0.340745
         6    -0.163484
         7    -0.229673
         8    -0.133273
         9    -0.142336
         dtype: float64
```

```
In [39]: test_data_values = test_data.values.tolist()
```

```
In [40]: likelihood_values = likelihood.values.tolist()
```

```

In [41]: #Question 3
# -1 mark: wrong answer, correct approach
# -2 marks: wrong approach, no answer
pred = []
for i in range(0,1000):
    prob_list = []
    for cls in range(0,10):
        prob_val = np.log(prior.loc[cls])
        for j in range(0,784):
            prob_val += test_data_values[i]
[j]*np.log(likelihood_values[cls][j]) + (1-test_data_values[i][j])*np.log(1-likelihood_values[cls][j])
        prob_list.append(prob_val)
    pred.append(np.argmax(prob_list))

```

```

In [42]: accuracy_score(test_labels, pred)

```

```

Out[42]: 0.809000000000000005

```

```

In [43]: #Question 4
# -1 mark: wrong answer, correct approach
# -2 marks: wrong approach, no answer

from sklearn.preprocessing import normalize
cnf_matrix = confusion_matrix(test_labels, pred)
print (cnf_matrix)
#cnf_matrix_norm = normalize(cnf_matrix,axis=1, norm='l1')
#print (np.around(cnf_matrix_norm,2))
#Answer: 4-9 (21),5-3 (11), 8-3 (9)

```


```

[[ 74  0  0  0  0  5  2  0  4  0]
 [ 0 120  0  0  0  4  1  0  1  0]
 [ 1  7 88  4  0  1  2  3  8  2]
 [ 0  2  1 86  1  6  3  2  3  3]
 [ 1  1  1  0 83  0  2  0  1 21]
 [ 3  1  1 11  2 62  2  3  1  1]
 [ 3  0  4  0  3  4 73  0  0  0]
 [ 0  6  2  0  3  1  0 77  3  7]
 [ 0  2  2  9  4  3  1  2 61  5]
 [ 0  1  0  1  4  0  0  0  3 85]]

```


```
In [44]: plt.imshow(np.reshape(test_data.loc[8], (28,28)))
```

```
Out[44]: <matplotlib.image.AxesImage at 0x11a8f3ed0>
```



```
In [45]: plt.imshow(np.reshape(test_data.loc[15], (28,28)))
```

```
Out[45]: <matplotlib.image.AxesImage at 0x11a9849d0>
```



GMM and LDA

```
In [46]: from sklearn import datasets

bc_data = datasets.load_breast_cancer()

n = bc_data.data.shape[0]

train_data = bc_data.data[:int(0.4*n)]
val_data = bc_data.data[int(0.4*n):int(0.6*n)]
test_data = bc_data.data[int(0.6*n):]
train_labels = bc_data.target[:int(0.4*n)]
val_labels = bc_data.target[int(0.4*n):int(0.6*n)]
test_labels = bc_data.target[int(0.6*n):]

for cov in ['full', 'tied', 'diag', 'spherical']:
    estimator = GaussianMixture(n_components=2,
                                covariance_type=cov, random_state=0)

    estimator.means_init = np.array([train_data[train_labels ==
i].mean(axis=0)
                                for i in range(2)])
    estimator.fit(train_data, train_labels)
    pred = estimator.predict(val_data)
    print ('Validation accuracy for covariance type ' + cov + ' = ' +
str(accuracy_score(val_labels, pred)))
```

```
Validation accuracy for covariance type full = 0.912280701754
Validation accuracy for covariance type tied = 0.859649122807
Validation accuracy for covariance type diag = 0.947368421053
Validation accuracy for covariance type spherical = 0.973684210526
```

```
In [47]: # Best for diag

train_data = np.concatenate((train_data, val_data))
train_labels = np.concatenate((train_labels, val_labels))

estimator = GaussianMixture(n_components=2,
                             covariance_type='spherical', random_state=0)

estimator.means_init = np.array([train_data[train_labels ==
i].mean(axis=0)
                             for i in range(2)])
estimator.fit(train_data, train_labels)
pred = estimator.predict(test_data)
print ('Test accuracy for covariance type ' + cov + ' = ' + str(accuracy_
score(test_labels, pred)))
```

```
Test accuracy for covariance type spherical = 0.938596491228
```

```
In [48]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

# Intialize
clf = LinearDiscriminantAnalysis()
# Train
clf.fit(train_data, train_labels)
# Test
pred = clf.predict(test_data)

print ('Test accuracy for covariance type ' + cov + ' = ' + str(accuracy_
score(test_labels, pred)))

Test accuracy for covariance type spherical = 0.973684210526
```

```
In [55]: clf.coef_
```

```
Out[55]: array([[ -0.00539085, -0.00539085, -0.00539085, -0.00539085, -0.0053908
5,
               -0.00539085, -0.03836087, -0.03836087, -0.00539085, -0.0053908
5,
               -0.00539085, -0.00539085, -0.00539085, -0.00539085, -0.0053908
5,
               -0.00539085, -0.03836087, -0.03836087, -0.00539085, -0.0053908
5,
               -0.00539085, -0.00539085, -0.00539085, -0.00539085, -0.0053908
5,
               -0.00539085, -0.03836087, -0.03836087, -0.00539085, -0.0053908
5]])
```

```
In [56]: clf.intercept_
```

```
Out[56]: array([-0.61694672])
```

```
In [73]: #Crosscheck
from sklearn.naive_bayes import BernoulliNB
clf = BernoulliNB()
clf.fit(train_data, train_labels)
```

```
Out[73]: BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True)
```

```
In [74]: pred = clf.predict(test_data)
```

```
In [75]: accuracy_score(test_labels, pred)
```

```
Out[75]: 0.809000000000000005
```

```
In [76]: clf.feature_log_prob_
```

```
Out[76]: array([[ -6.38687932,  -6.38687932,  -6.38687932, ...,  -6.38687932,
                -6.38687932,  -6.38687932],
               [ -6.51174533,  -6.51174533,  -6.51174533, ...,  -6.51174533,
                -6.51174533,  -6.51174533],
               [ -6.36818719,  -6.36818719,  -6.36818719, ...,  -6.36818719,
                -6.36818719,  -6.36818719],
               ...,
               [ -6.48157713,  -6.48157713,  -6.48157713, ...,  -6.48157713,
                -6.48157713,  -6.48157713],
               [ -6.315358   ,  -6.315358   ,  -6.315358   , ...,  -6.315358   ,
                -6.315358   ,  -6.315358   ],
               [ -6.4019172  ,  -6.4019172  ,  -6.4019172  , ...,  -6.4019172  ,
                -6.4019172  ,  -6.4019172  ]])
```