

---

# *Unsupervised Learning with Clustering*

**K-means**

**EM clustering**

**Hierarchical Clustering**

---

# *Why Clustering*


- **A good grouping implies some structure**
- **In other words, given a good grouping, we can then:**
  - Interpret and label clusters
  - Identify important features
  - Characterize new points by the closest cluster (or nearest neighbors)
  - Use the cluster assignments as a compression or summary of the data

---

# *Clustering*

- **Basic idea: Group similar things together**
- **Unsupervised Learning – Useful when no other info is available**
- **K-means**
  - Partitioning instances into k disjoint clusters
  - Measure of similarity

# Clustering



---

# *Clustering Techniques*

- **K-means clustering**
- **Hierarchical clustering**
- **Conceptual clustering**
- **Probability-based clustering**
- **Bayesian clustering**

---

# *Clustering Techniques*

- **K-means clustering**
- **Hierarchical clustering**
- **Conceptual clustering**
- **Probability-based clustering**
- **Bayesian clustering**

---

# *Clustering Applications*

- **Example: Clustering of a large number of gene experiments**
- **Multiple sequence alignment of genes closely clustered together**
- **Search for metabolic pathways genes may be involved with**
- **Possible functional classification of genes in the same cluster**
- **Identifying co-regulated genes from expression arrays**

---

# ***Common uses of Clustering***

- **Often used as an exploratory data analysis tool**
- **In one-dimension, a good way to quantify real-valued variables into  $k$  non-uniform buckets**
- **Used on acoustic data in speech understanding to convert waveforms into one of  $k$  categories (known as Vector Quantization)**
- **Also used for choosing color palettes on old fashioned graphical display devices**
- **Color Image Segmentation**



---


# *Clustering*

- **Unsupervised: no target value to be predicted**
- **Differences ways clustering results can be produced/represented/learned**
  - Exclusive vs. overlapping
  - Deterministic vs. probabilistic
  - Hierarchical vs. flat
  - Incremental vs. batch learning

# *Clustering Objective*

- **Objective: find subsets that are similar within cluster and dissimilar between clusters**
- **Similarity defined by distance measures**

- Euclidean distance
- Manhattan distance
- Mahalanobis  
(Euclidean w/dimensions  
rescaled by variance)



# Clustering Objective

- **Objective: find subsets that are similar within cluster and dissimilar between clusters**
- **Similarity defined by distance measures**



- Euclidean distance =  
 $\text{sqrt}[(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots]$

- Manhattan distance  
 $[|a_1 - b_1| + |a_2 - b_2| + \dots]$

- Cosine (insensitive to size)

Euc Dist/

$$\text{sqrt}[(a_1)^2 + (a_2)^2 + \dots] \cdot \text{sqrt}[(b_1)^2 + \dots]$$



---

# ***The $k$ -means Algorithm***

## ***Iterative Distance Based Clustering***

- **Clusters the data into  $k$  groups where  $k$  is specified in advance**
  1. Cluster centers are chosen at random
  2. Instances are assigned to clusters based on their distance to the cluster centers
  3. Centroids of clusters are computed – “means”
  4. Go to 1st step until convergence

# ***K-means Clustering***

- **A simple, effective, and standard method**

Start with K initial cluster centers

Loop:

Assign each data point to nearest cluster center

Calculate mean of cluster for new center

Stop when assignments don't change

- **Issues:**

How to choose K?

How to choose initial centers?

Will it always stop?


---

# ***K-Means Clustering Pros & Cons***

- **Simple and reasonably effective**
- **The final cluster centers do not represent a global minimum but only a local one**
- **Result can vary significantly based on initial choice of seeds**
  - Completely different final clusters can arise from differences in the initial randomly chosen cluster centers
- **Algorithm can easily fail to find a reasonable clustering**

# *Getting Trapped in a Local Minimum*

- **Example: four instances at the vertices of a two-dimensional rectangle**
  - Local minimum: two cluster centers at the midpoints of the rectangle's long sides
- **Simple way to increase chance of finding a global optimum: restart with different random seeds**



---

# Clustering

- Partition unlabeled examples into disjoint subsets of **clusters**, such that:
  - Examples within a cluster are very similar
  - Examples in different clusters are very different
- Discover new categories in an **unsupervised** manner (no sample category labels provided)



# *K-Means Algorithm*

Let  $d$  be the distance measure between instances.

Select  $k$  random instances  $\{s_1, s_2, \dots, s_k\}$  as seeds.

Until clustering converges or other stopping criterion:

For each instance  $x_i$ :


Assign  $x_i$  to the cluster  $c_j$  such that  $d(x_i, s_j)$  is minimal.

*(Update the seeds to the centroid of each cluster)*

For each cluster  $c_j$

$$s_j = \mu(c_j)$$

# *K Means Example* ( $K=2$ )



Pick seeds

Reassign clusters

Compute centroids

Reassign clusters

Compute centroids

Reassign clusters

**Converged!**


---

# *Seed Choice*

- **Results can vary based on random seed selection**
- **Some seeds can result in poor convergence rate, or convergence to sub-optimal clusters**
- **Select good seeds using a heuristic or the results of another method**


# *K-means Example*

- For  $K=1$ , using Euclidean distance, where will the cluster center be?



# *K-means Example*

- For  $K=1$ , the overall mean minimizes Sum Squared Error (SSE), aka Euclidean distance







Simple example:

#choose 1 data point as initial K centers





#10 is max loop iterations

#1 is number of initial sets to try

# *K-means Example*




# K-means Example



As K increases  
individual points  
get a cluster

# Choosing $K$ for $K$ -means




- Not much improvement after  $K=2$  (“elbow”)




# *K-means Example – more points*

*How many clusters should there be?*



# Choosing $K$ for $K$ -means





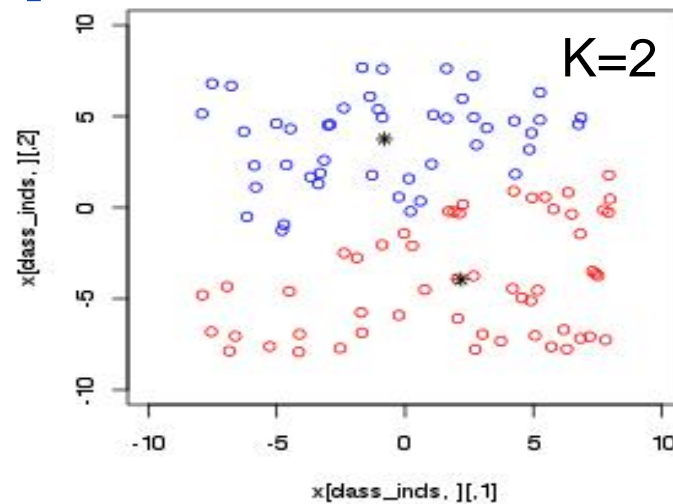

- Smooth decrease at  $K \geq 2$ , harder to choose
- In general, smoother decrease  $\Rightarrow$  less structure

---


# ***K-means Guidelines***

- **Choosing K:**
  - “Elbow” in total-within-cluster SSE as  $K=1 \dots N$
  - Cross-validation: hold out points, compare fit as  $K=1 \dots N$
- **Choosing initial starting points:**
  - take  $K$  random data points, do several K-means, take best fit
- **Stopping:**
  - may converge to sub-optimal clusters
  - may get stuck or have slow convergence (point assignments bounce around), 10 iterations is often good

# *K-means Example: uniform dist.*



# Choosing $K$ - uniform



- Smooth decrease across  $K \Rightarrow$  less structure

---

# *K-means Clustering Issues*

- **Scale:**
  - Dimensions with large numbers may dominate distance metrics
- **Outliers:**
  - Outliers can pull cluster mean, K-medoids uses median instead of mean

# *Probability-based Clustering*

- **Problems with K-means & Hierarchical methods:**
  - Division by k
  - Order of examples
  - Merging/splitting operations might not be sufficient to reverse the effects of bad initial ordering
  - Is result at least local minimum of category utility?
- **Solution:**
  - Find the most likely clusters given the data
- **Instance has certain probability of belonging to a particular cluster**

# Soft Clustering

- So far clustering methods assumed that each instance has a “hard” assignment to exactly one cluster
- No uncertainty about class membership or an instance belonging to more than one cluster
- *Soft clustering* gives probabilities that an instance belongs to each of a set of clusters
- Each instance is assigned a probability distribution across a set of discovered clusters
  - probabilities of all categories must sum to 1



---

# *Soft Clustering Methods*

- **Fuzzy Clustering**

- Use weighted assignments to all clusters
- Weights depend on relative distance
- Find min weighted SSE

- **Expectation-Maximization:**


- Initialize a mixture of multivariate Gaussian distributions
- Find means, variances, and mixture weights that maximize probability of data

# *Finite mixtures*

- Probabilistic clustering algorithms model the data using a mixture of distributions
- Each cluster is represented by one distribution
  - The distribution governs the probabilities of attributes values in the corresponding cluster
- They are called finite mixtures because there is only a finite number of clusters being represented
- Usually individual distributions are normal
- Distributions are combined using cluster weights

# A Two-Class Mixture Model

data											
A	51	B	62	B	64	A	48	A	39	A	51
A	43	A	47	A	51	B	64	B	62	A	48
B	62	A	52	A	52	A	51	B	64	B	64
B	64	B	64	B	62	B	63	A	52	A	42
A	45	A	51	A	49	A	43	B	63	A	48
A	42	B	65	A	48	B	65	B	64	A	41
A	46	A	48	B	62	B	66	A	48		
A	45	A	49	A	43	B	65	B	64		
A	45	A	46	A	40	A	46	A	48		



## *Using the Mixture Model*

- The probability of an instance  $x$  belonging to cluster  $A$  is:

$$PR[A|x] = \frac{Pr[x|A]Pr[A]}{Pr[x]} = \frac{f(x; \mu_A, \sigma_A)p_A}{Pr[x]}$$

# *Learning the Clusters*

- Assume we know that there are  $k$  clusters
- To learn the clusters we need to determine their parameters
  - I.e. their means and standard deviations
- Start with the initial guess for the 5 parameters use them to calculate cluster probabilities for each instance, use these probabilities to re estimate the parameters and repeat
- We actually have a performance criterion: the likelihood of the training data given the clusters

# *Expectation Maximization (EM)*

- Probabilistic method for soft clustering
- Iterative method for learning probabilistic categorization model from unsupervised data
- Direct method that assumes  $k$  clusters:  $\{c_1, c_2, \dots, c_k\}$
- Soft version of  $k$ -means
- Assumes a probabilistic model of categories that allows computing  $P(c_i | E)$  for each category,  $c_i$ , for a given example,  $E$

---

# *EM Algorithm*


- Initially assume random assignment of examples to categories
- Learn an initial probabilistic model by estimating model parameters from this randomly labeled data

# *The EM Algorithm*

- **EM algorithm:**
  - expectation-maximization algorithm
- **Generalization of k-means to probabilistic setting**
- **Similar iterative procedure**
  1. Calculate cluster probability for each instance (expectation step)
  2. Estimate distribution parameters based on the cluster probabilities (maximization step)
- **Cluster probabilities are stored as instance weights**







# *Kmeans – unequal cluster variance*




Can you guess K?

# *K-means – unequal cluster variance*




# Choosing $K$ – unequal distributions



- Smooth decrease across  $K \Rightarrow$  less structure

# *EM clustering*

**Classification**



- Selects  $K=2$   
(either by Information Criterion=  
min of  $SSE + K \cdot \log N$ ,  
Or by cross-validation)
- Handles unequal variance

---


# ***K-means computations***

- **Distance of each point to each cluster center**
  - For  $N$  points,  $D$  dimensions: each loop requires  $N*D*K$  operations
- **Update Cluster centers**
  - only track points that change, get change in cluster center
- **But for EM errors to each cluster center update a probability function**

# *K-means vs EM performance*

1 Gordon compute node, normal random matrices

R: `system.time(Mclust())`




# *Kmeans big data example*


- **45,000 NYTimes articles, 102,000 unique words**

(UCI Machine Learning repository)

- **Full Data Matrix: 45Kx102K ~ 40Gb**



## *Kmeans results*




## 7 viable clusters found



## Other distance measures

- Cosine: each row is treated as vector in  $R^p$ , then take angles



- Jaccard (over sets A,B):  $1 - (|A \cap B| / |A \cup B|)$

## *Other distance measures*

- Hamming distance: count 1 if values different  
e.g. appropriate for binary strings

010100110  
000100010

↑                      ↑

Total difference is 2

---

# *Summary*

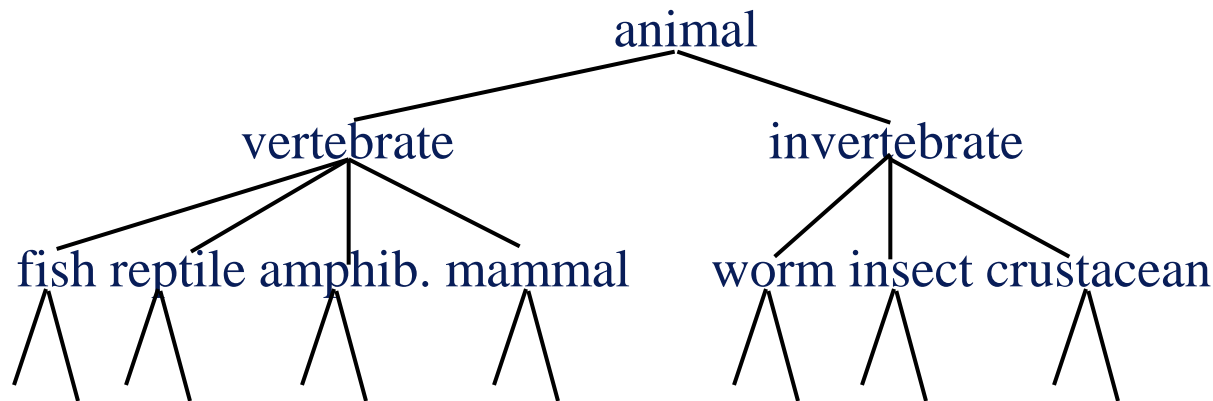
- **Labeled clusters can be interpreted by using supervised learning - train a tree or learn rules**
- **Can be used to fill in missing attribute values**
- **All methods have a basic assumption of independence between the attributes**
  - Some methods allow the user to specify in advanced that two of more attributes are dependent and should be modeled with a joint probability

# Clustering algorithms in scikit-learn

Method name	Parameters	Scalability	Usecase	Geometry (metric used)
<u><a href="#">K-Means</a></u>	number of clusters	Very large n_samples, medium n_clusters with <u><a href="#">MiniBatch code</a></u>	General-purpose, even cluster size, flat geometry, not too many clusters	Distances between points
<u><a href="#">Affinity propagation</a></u>	damping, sample preference	Not scalable with n_samples	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
<u><a href="#">Mean-shift</a></u>	bandwidth	Not scalable with n_samples	Many clusters, uneven cluster size, non-flat geometry	Distances between points
<u><a href="#">Spectral clustering</a></u>	number of clusters	Medium n_samples, small n_clusters	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
<u><a href="#">Ward hierarchical clustering</a></u>	number of clusters	Large n_samples and n_clusters	Many clusters, possibly connectivity constraints	Distances between points
<u><a href="#">Agglomerative clustering</a></u>	number of clusters, linkage type, distance	Large n_samples and n_clusters	Many clusters, possibly connectivity constraints, non Euclidean distances	Any pairwise distance
<u><a href="#">DBSCAN</a></u>	neighborhood size	Very large n_samples, medium n_clusters	Non-flat geometry, uneven cluster sizes	Distances between nearest points
<u><a href="#">Gaussian mixtures</a></u>	many	Not scalable	Flat geometry, good for density estimation	Mahalanobis distances to centers
<u><a href="#">Birch</a></u>	branching factor, threshold, optional global clusterer.	Large n_clusters and n_samples	Large dataset, outlier removal, data reduction.	Euclidean distance between points

# *Hierarchical Clustering*

- Build a tree-based hierarchical taxonomy (*dendrogram*) from a set of unlabeled examples



- Recursive application of a standard clustering algorithm can produce a hierarchical clustering

---

# ***Incremental & Hierarchical Clustering***

- **Start with 1 cluster (all instances) and do splits**  
**OR**  
**Start with N clusters (1 per instance) and do merges**
- **Can be greedy & expensive in its search**  
some algorithms might merge & split  
algorithms need to store and recalculate distances
- **Need distance between groups**  
in contrast to K-means

---

# ***Incremental & Hierarchical Clustering***

- **Result is a hierarchy of clusters**
  - displayed as a 'dendrogram' tree
- **Useful for tree-like interpretations**
  - syntax (e.g. word co-occurrences)
  - concepts (e.g. classification of animals)
  - topics (e.g. sorting Enron emails)
  - spatial data (e.g. city distances)
  - genetic expression (e.g. possible biological networks)
  - exploratory analysis

---

# *Incremental Clustering*




- **Works incrementally instance by instance forming a hierarchy of clusters**
- **COBWEB - nominal; CLASSIT– numeric attributes**
- **Instances are added one at the time**
  - Tree is updated appropriately at each step
  - Finding the right leaf for an instance
  - Restructuring the tree
- **How and where to update based on category utility value**



# Clustering: Weather Data


Weather Data Set				
ID Code	Outlook	Temperature	Humidity	Windy
A	Sunny	Hot	High	False
B	Sunny	Hot	High	True
C	Overcast	Hot	High	False
D	Rain	Mild	High	False
E	Rain	Cool	Normal	False
F	Rain	Cool	Normal	True
F	Overcast	Cool	Normal	True
H	Sunny	Mild	High	False
I	Sunny	Cool	Normal	False
J	Rain	Mild	Normal	False
K	Sunny	Mild	Normal	True
L	Overcast	Mild	High	True
M	Overcast	Hot	Normal	False
N	Rain	Mild	High	True

# Clustering




# Clustering

Step 4



Step 5




---

# *Merging*

- **Consider all pairs of nodes for merging and evaluate category utility of each**
  - Computationally expensive
- **When scanning nodes for a suitable host – both the best matching node and the runner-up are noted**
- **The best will form the host for new instance unless merging host and runner-up produces better CU**

# *Final Hierarchy*



\_\_\_\_\_



# Category utility

- **Category utility is a kind of quadratic loss function defined on conditional probabilities:**

$$CU(C_1, C_2, \dots, C_k) = \frac{\sum_l \Pr[C_l] \sum_i \sum_j (\Pr[a_i = v_{ij} | C_l]^2 - \Pr[a_i = v_{ij}]^2)}{k}$$

- **$C_1, \dots, C_k$  are  $k$  clusters**
- **$a_i$  is the  $i$ th attribute**
- **Takes on values  $v_{i1}, v_{i2}, \dots$**

# *Category Utility Extended to Numeric Attributes*

- **Assuming normal distribution:**

$$CU = \frac{\sum_l \Pr[C_l] \frac{1}{2\sqrt{\pi}} \sum_i \left( \frac{1}{\sigma_{il}} - \frac{1}{\sigma_i} \right)}{k}$$

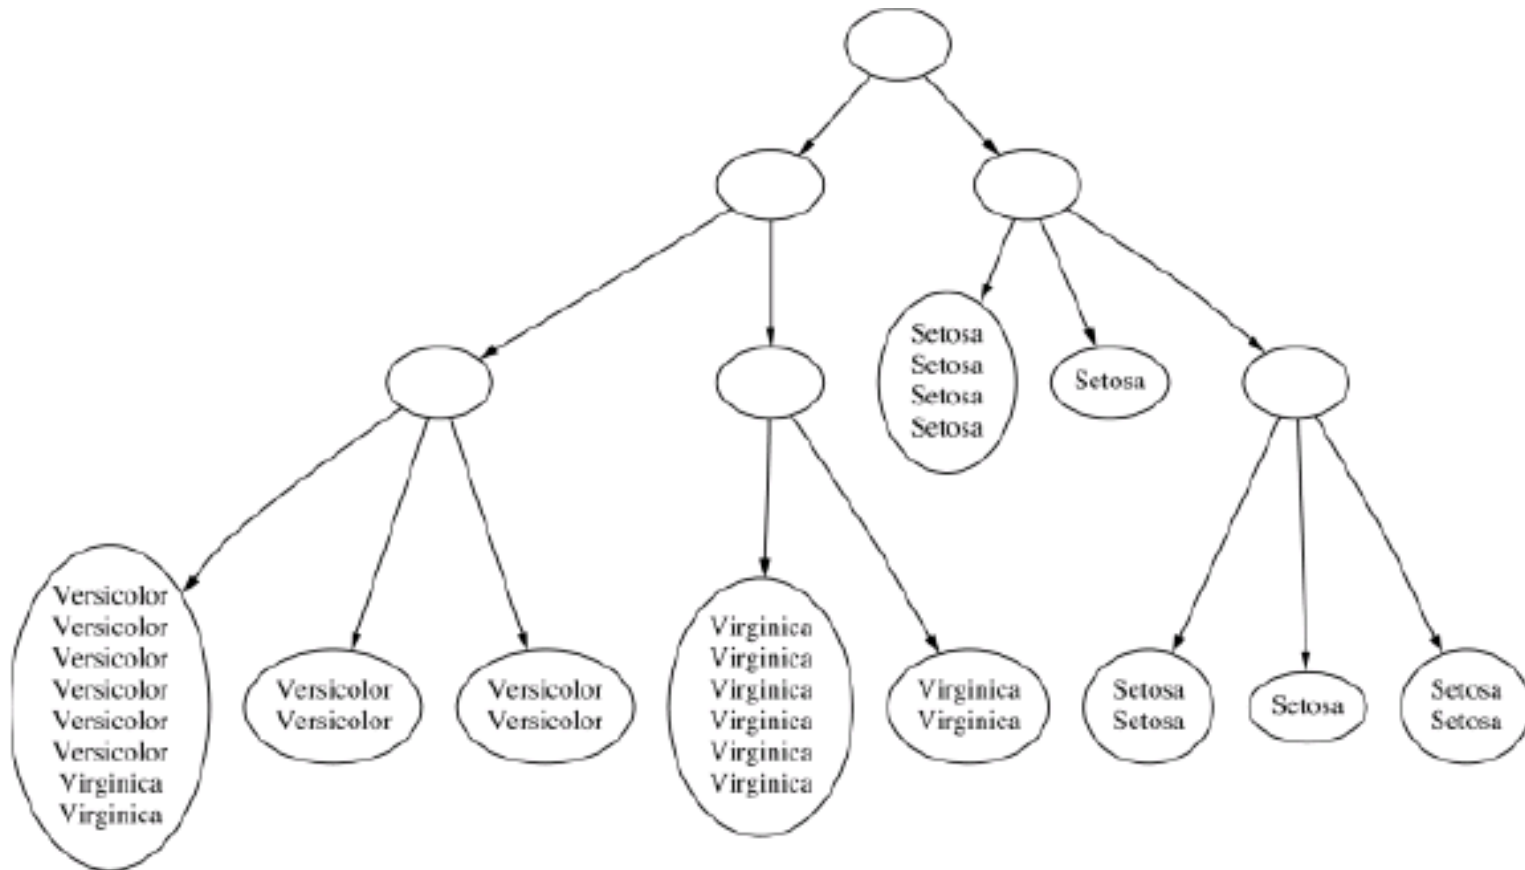
- **When Standard deviation of attribute  $a_i$  is zero it produced infinite value of the category utility formula**
- **Acuity parameter: pre-specified minimum variance on each attribute**
  - only one instance in a node produces 0 variance



# PACE



# Clustering with Cutoff




---

# ***Incremental & Hierarchical Clustering***


- **Clusters are merged/split according to distance or utility measure**
  - Euclidean distance (squared differences)
  - conditional probabilities (for nominal features)
- **Options to choose which clusters to ‘Link’**
  - single linkage, mean, average (w.r.t. points in clusters)  
(may lead to different trees, depending on spreads)
  - Ward method (smallest increase within cluster variance)
  - change in probability of features for given clusters

# *Linkage options*

- e.g. single linkage (closest to any cluster instance)




- e.g. mean (closest to mean of all cluster instances)




# Linkage options (cont')

- e.g. average (mean of pairwise distances)



- e.g. Ward's method (find new cluster with min. variance)



# ***Hierarchical Clustering Demo***


- **3888 Interactions among 685 proteins**

From Hu et.al. TAP dataset

<http://www.compsysbio.org/bacteriome/dataset/>)


b0009	b0014	0.92
b0009	b2231	0.87
b0014	b0169	1.0
b0014	b0595	0.76
b0014	b2614	1.0
b0014	b3339	0.95
b0014	b3636	0.9
b0015	b0014	0.99
.....		

# *Interactions as connections – structure is hard to see*



# *Hierarchical Clustering Demo*

- **hclust with “single” distance: chaining**






# *Hierarchical Clustering Demo*

- **hclust with “Ward” distance: spherical clusters**

Cluster Dendrogram




d2use

hclust ("", "ward")

# ***Hierarchical Clustering Demo***

- **Where height change looks big, cut off tree**

**Cluster Dendrogram**



---

# Summary

- **Having no label doesn't stop you from finding structure in data**
- **Labeled clusters can be interpreted by using supervised learning - train a tree or learn rules**
- **Can be used to fill in missing attribute values**
- **All methods have a basic assumption of independence between the attributes**
  - Some methods allow the user to specify in advanced that two or more attributes are dependent and should be modeled with a joint probability
- **Unsupervised methods are somewhat related**

---

## ***Exercise: Clustering Athlete's Data in Weka***

*Is there a relationship between athlete's physical attributes and their sport?*

Let's filter data to try a few, good, candidate sport categories.

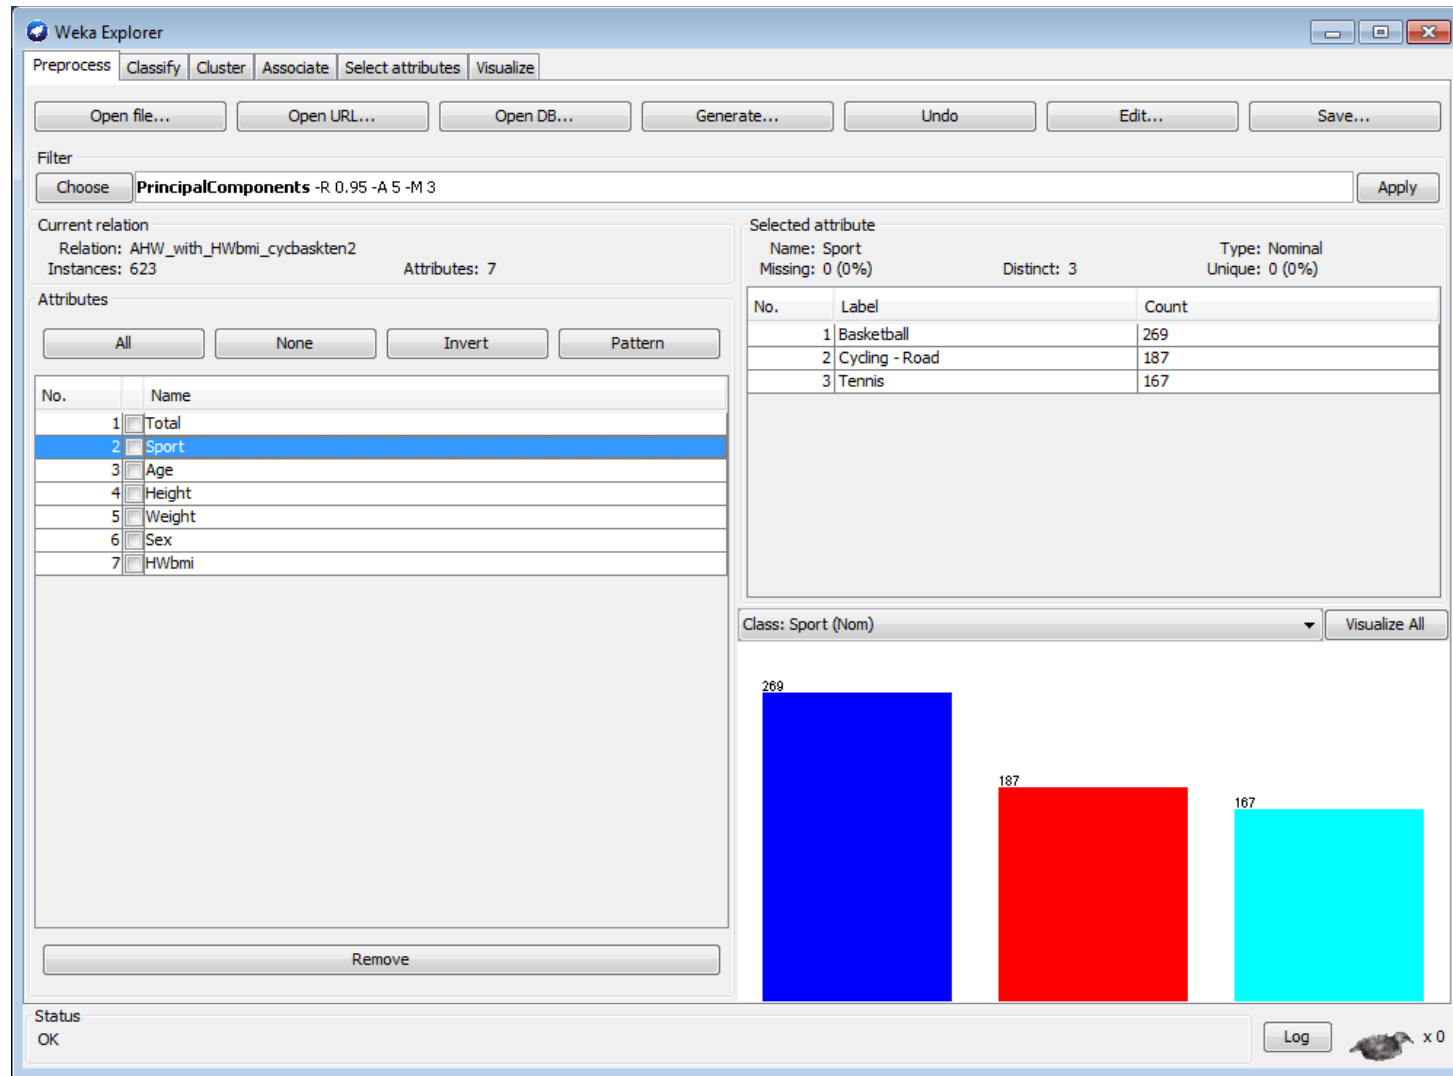
Weka not nice for picking out instances with multi-criteria, so I used excel:  
(Excel, data -> advanced -> click on column and then drop down selection menu.  
Include only a few sports)

Download AHW\_withHWbmi\_cycbaskten2.csv from [pace.sdsc.edu](http://pace.sdsc.edu), open in weka

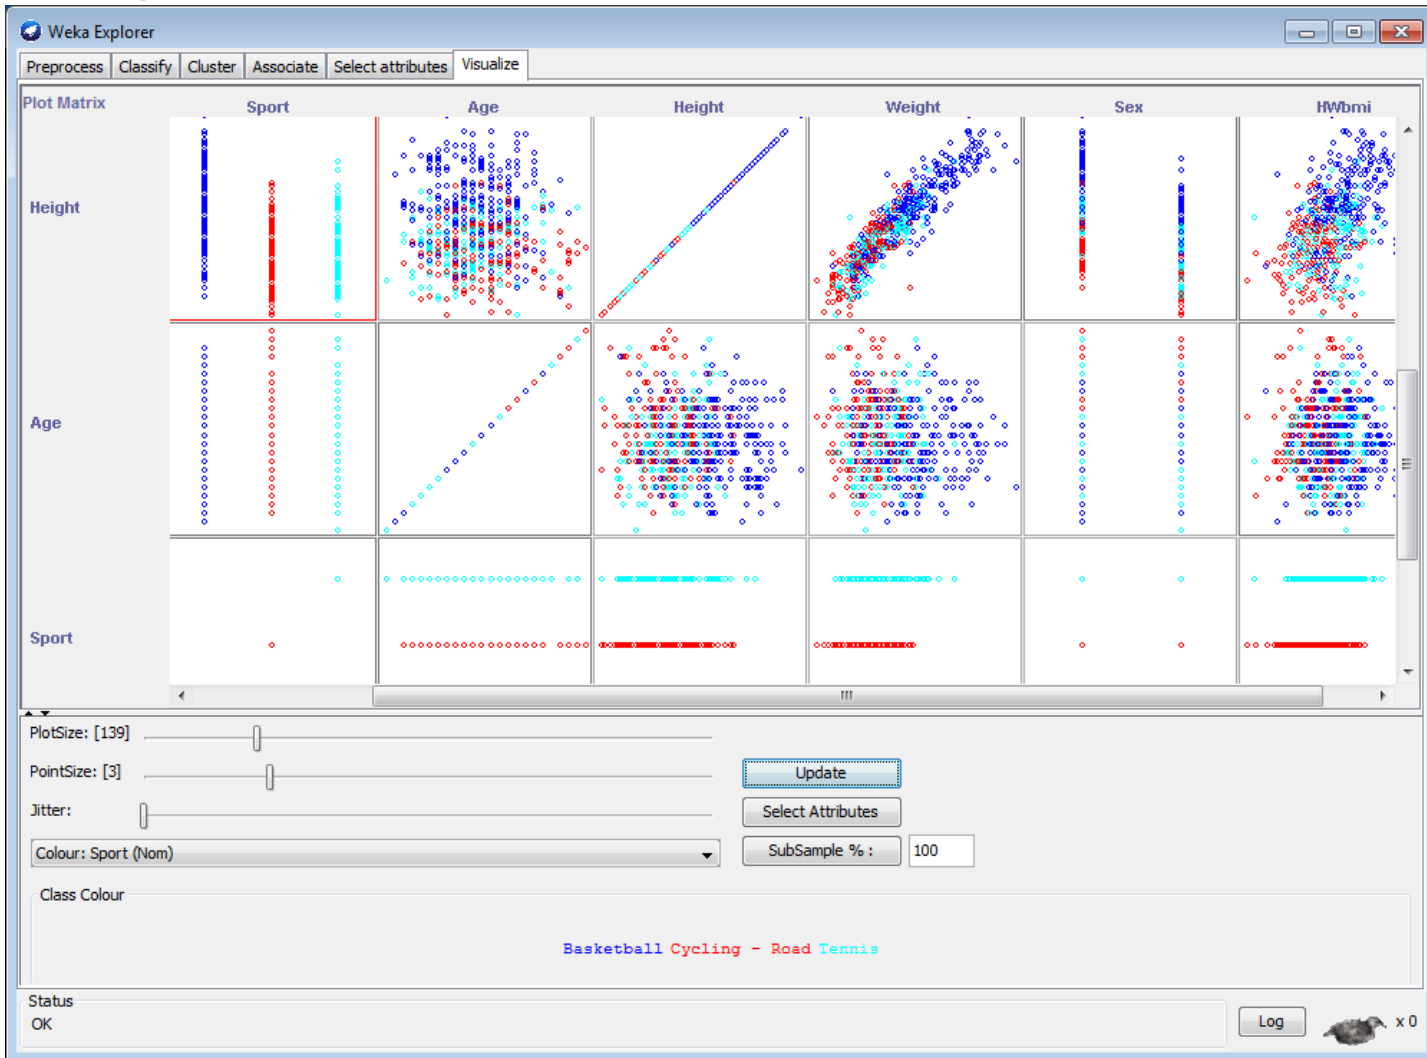
The screenshot shows a Microsoft Excel spreadsheet titled 'AHW\_1 - Microsoft Excel'. The 'Data' tab is active, and the 'Filter' button in the 'Sort & Filter' group is highlighted. A dropdown menu is open for the 'Sport' column, showing a list of sports with 'Basketball' selected. The data table has the following columns: Total, Sport, Age, Height, Weight, Sex. The data rows are as follows:

Total	Sport	Age	Height	Weight	Sex
23		170	60	M	
33		193	125	M	
30		187	76	M	
24				M	
26		178	85	F	
27		182	80	M	
30		182	73	F	
23		187	75	M	
27		190	80	M	
19		170		M	
37				M	
28				M	
28		173	60	F	
28		171	64	F	
22		171	62	M	
19		167		F	
22		175	63	M	
25		162	69	M	

# Notice the number of sports and their nominal values

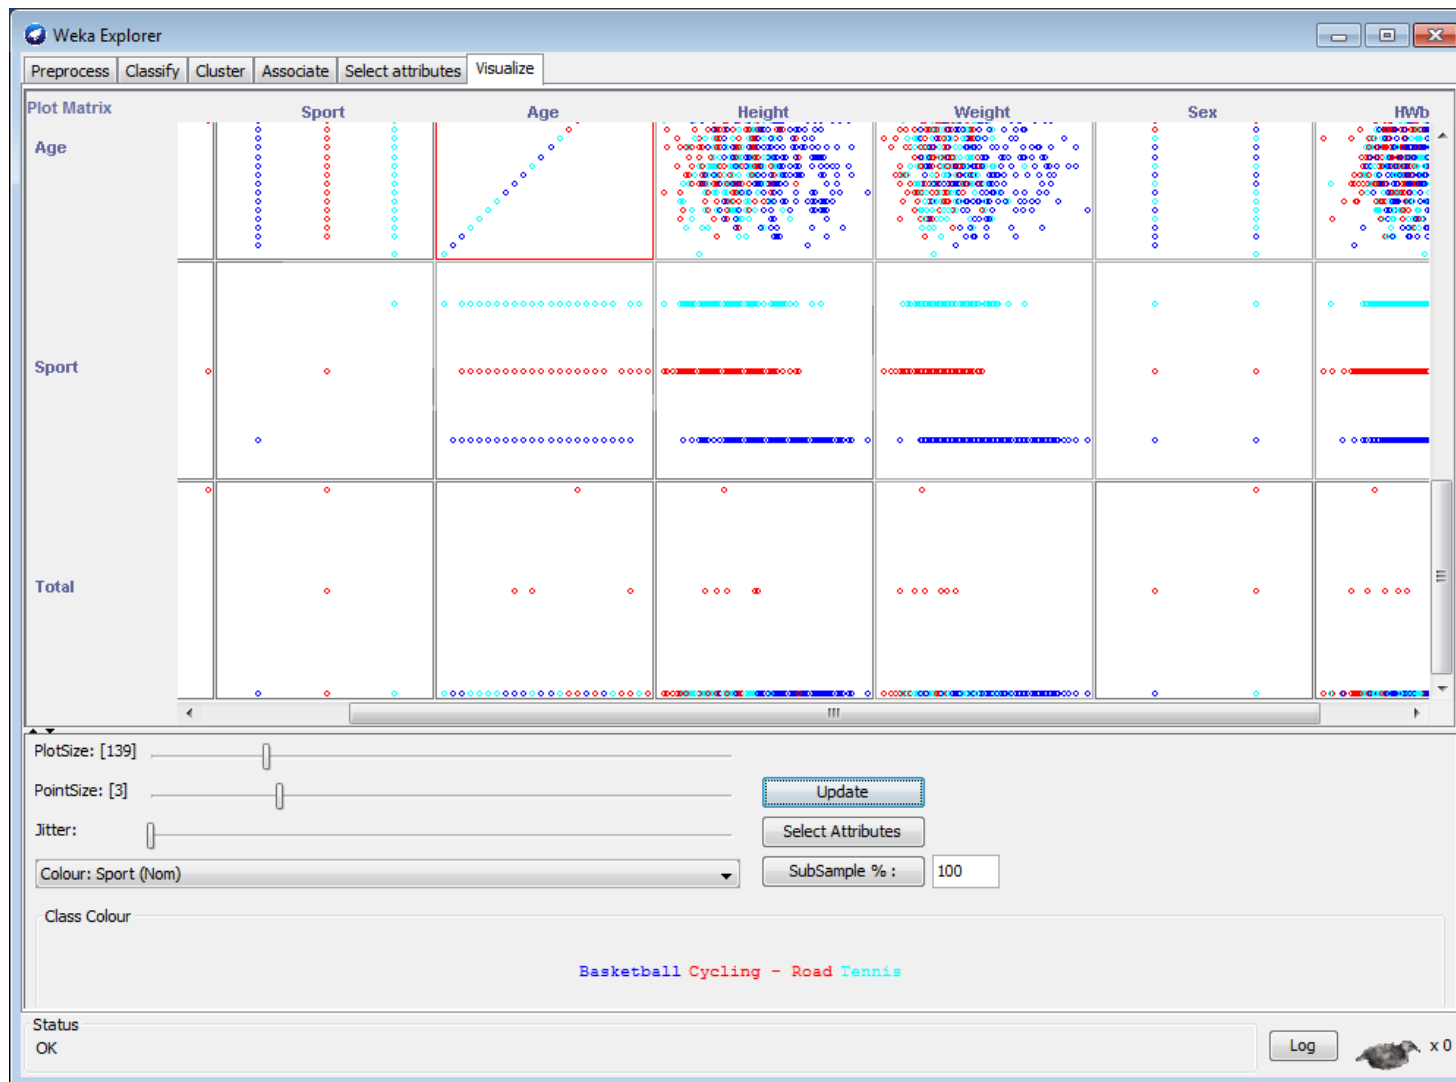


- Visualize correlation with sport as class – any problems, any promising relationships?



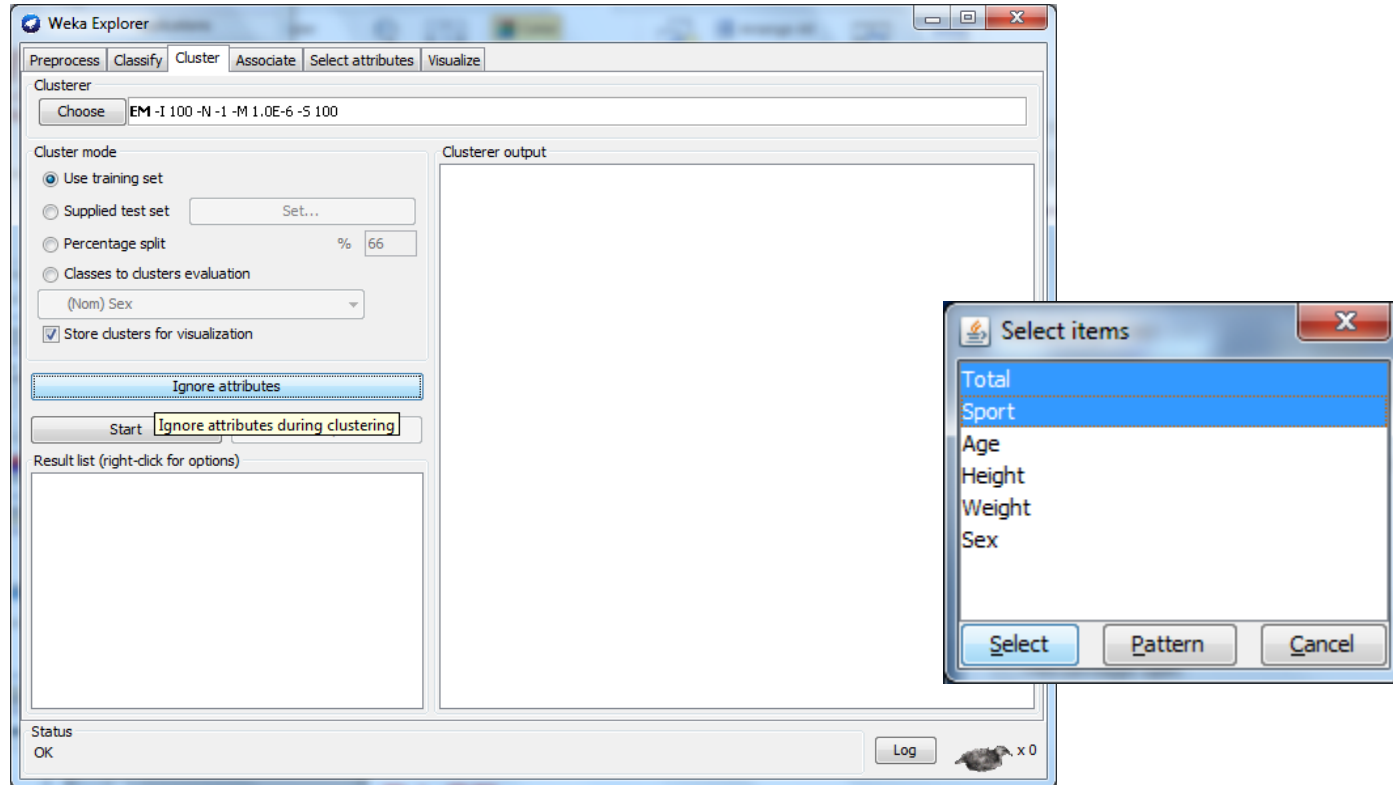
## Are there one or more variables we should ignore?

Predictive Analytics Center of Excellence

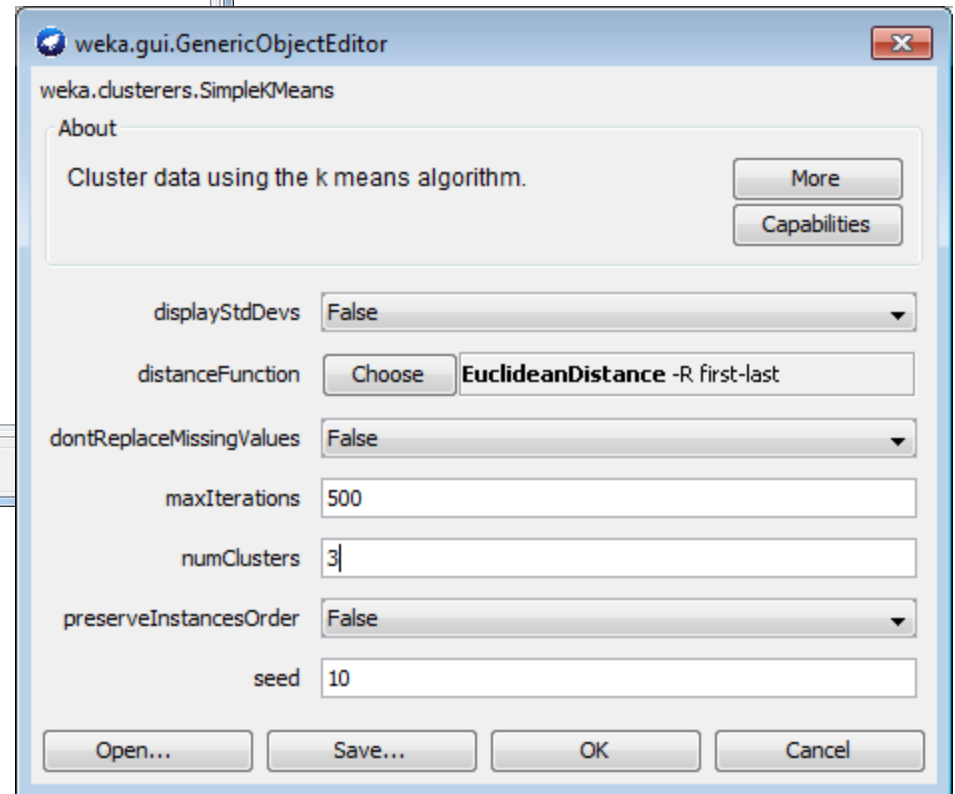





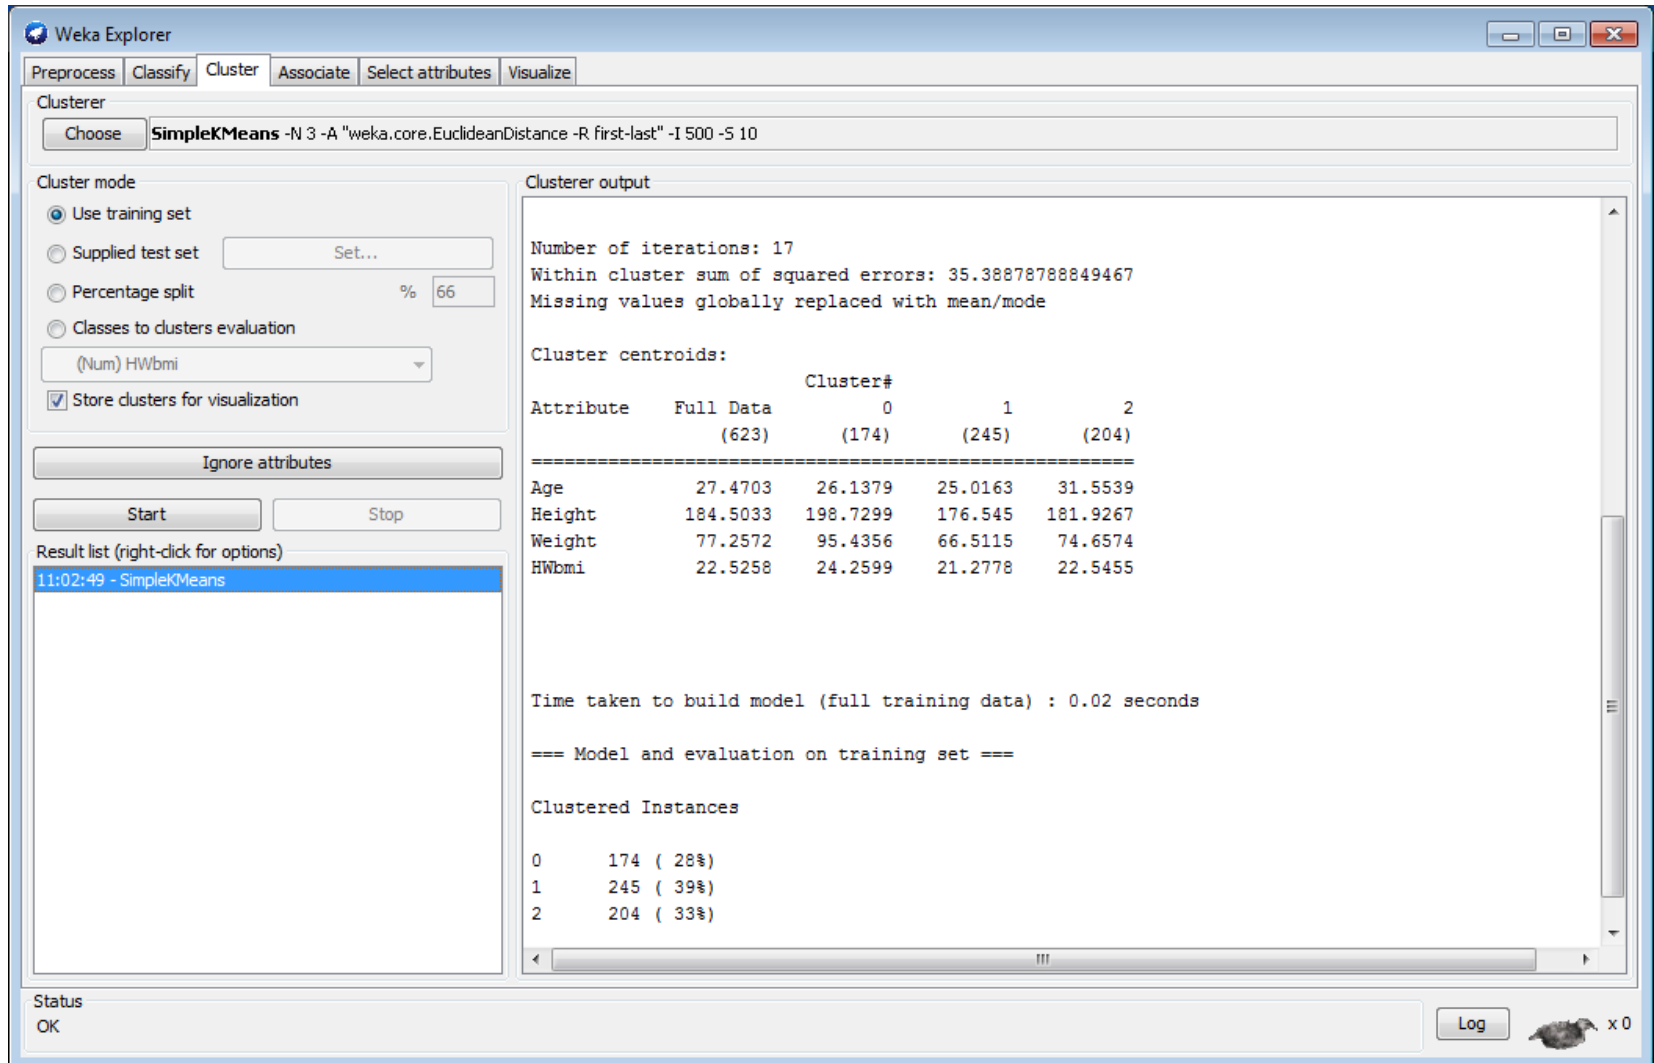
Select Ignore attributes -> ctrl-space to ignore sport (why?) and total medals won (why or why not?)  
What about sex attribute?



select cluster tab; choose -> simple Kmeans ; choose 3 classes, accept other defaults



Results end up in output panel. Take note of sum squared error. Any other message to note?



The screenshot shows the Weka Explorer application window. The 'Cluster' tab is selected. The 'Clusterer' dropdown is set to 'SimpleKMeans'. The 'Cluster mode' section has 'Use training set' selected. The 'Clusterer output' panel displays the following information:

Number of iterations: 17  
Within cluster sum of squared errors: 35.38878788849467  
Missing values globally replaced with mean/mode

Cluster centroids:

Attribute	Full Data (623)	Cluster# 0 (174)	1 (245)	2 (204)
Age	27.4703	26.1379	25.0163	31.5539
Height	184.5033	198.7299	176.545	181.9267
Weight	77.2572	95.4356	66.5115	74.6574
HWbmi	22.5258	24.2599	21.2778	22.5455

Time taken to build model (full training data) : 0.02 seconds

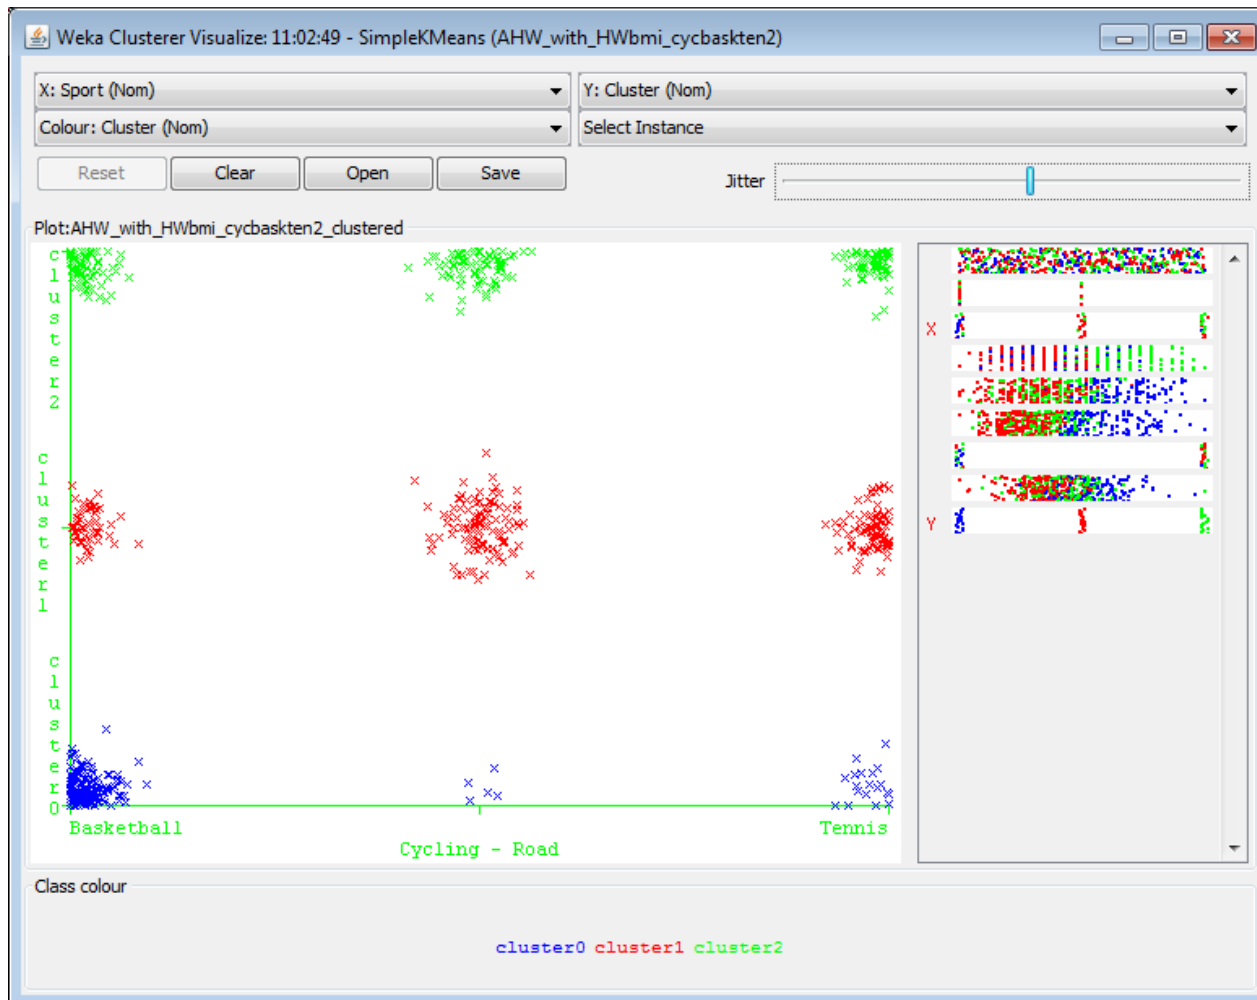
=== Model and evaluation on training set ===

Clustered Instances

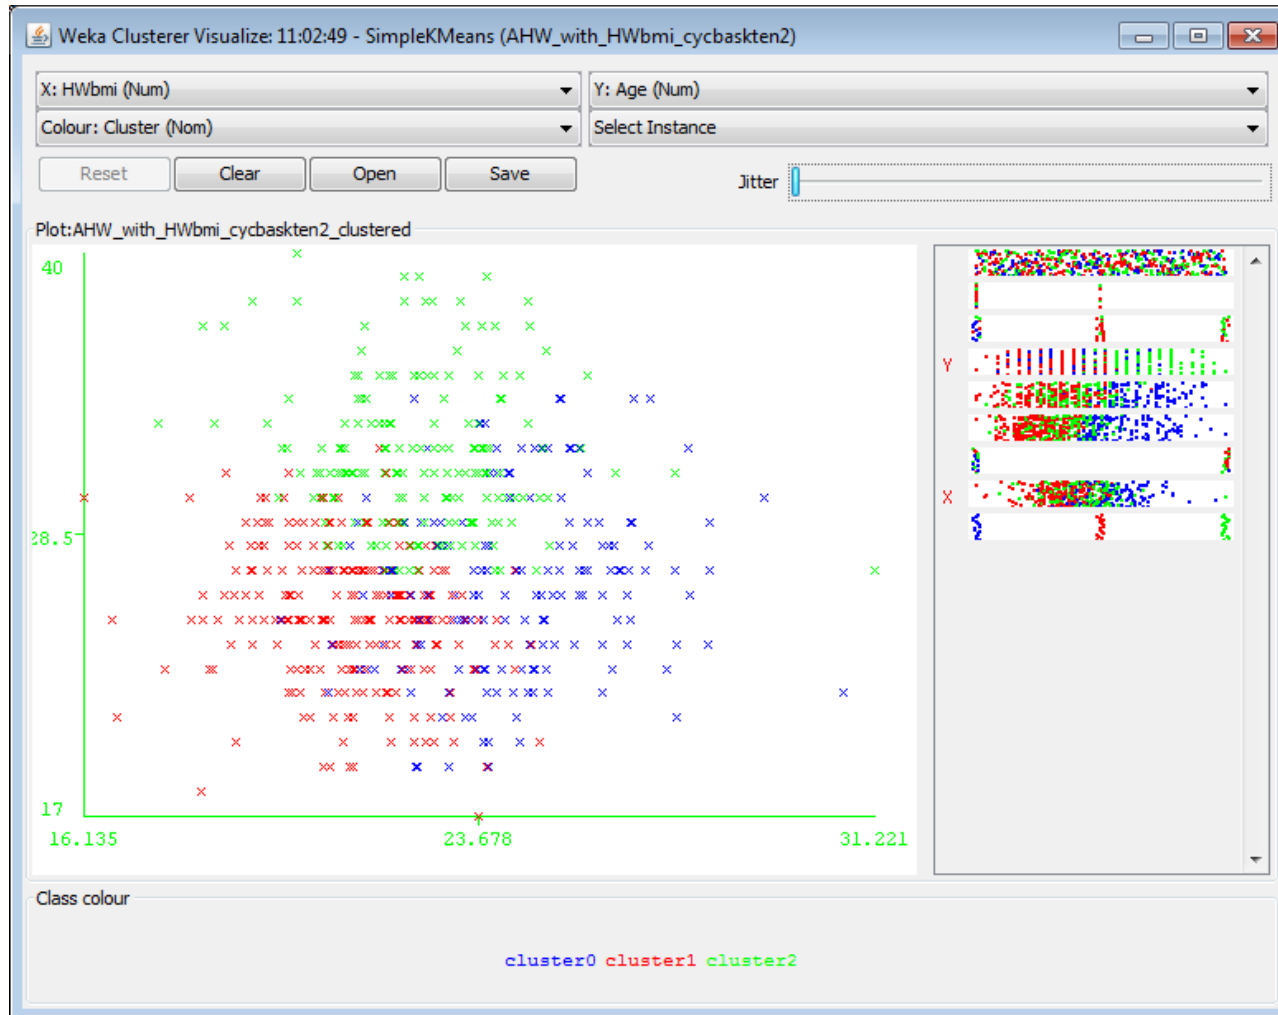
0	174 ( 28%)
1	245 ( 39%)
2	204 ( 33%)

The 'Result list' panel shows a single entry: '11:02:49 - SimpleKMeans'. The 'Status' bar at the bottom indicates 'OK'.

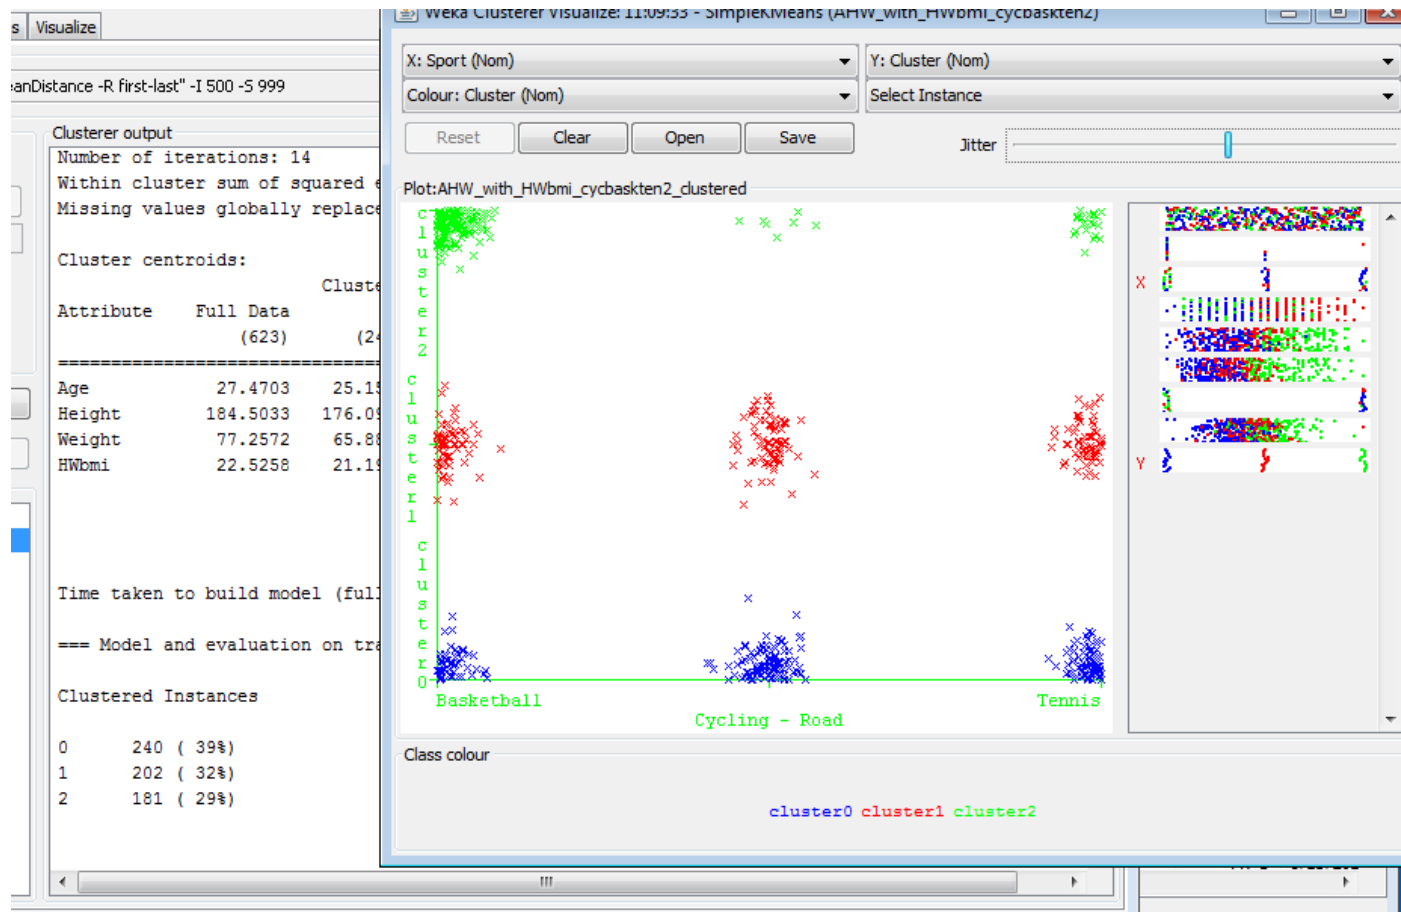
Right click on model result, visualize cluster assignment, add jitter  
Compare sport and cluster number – how do they correspond?  
Are the clusters useful to distinguish (or predict) an athlete's sport?



# Compare clusters in the 2 dimension subspace of BMI and Age - (or any other subspace)



Try rerunning with different random seed,  
you get different clustering, but similar SSE



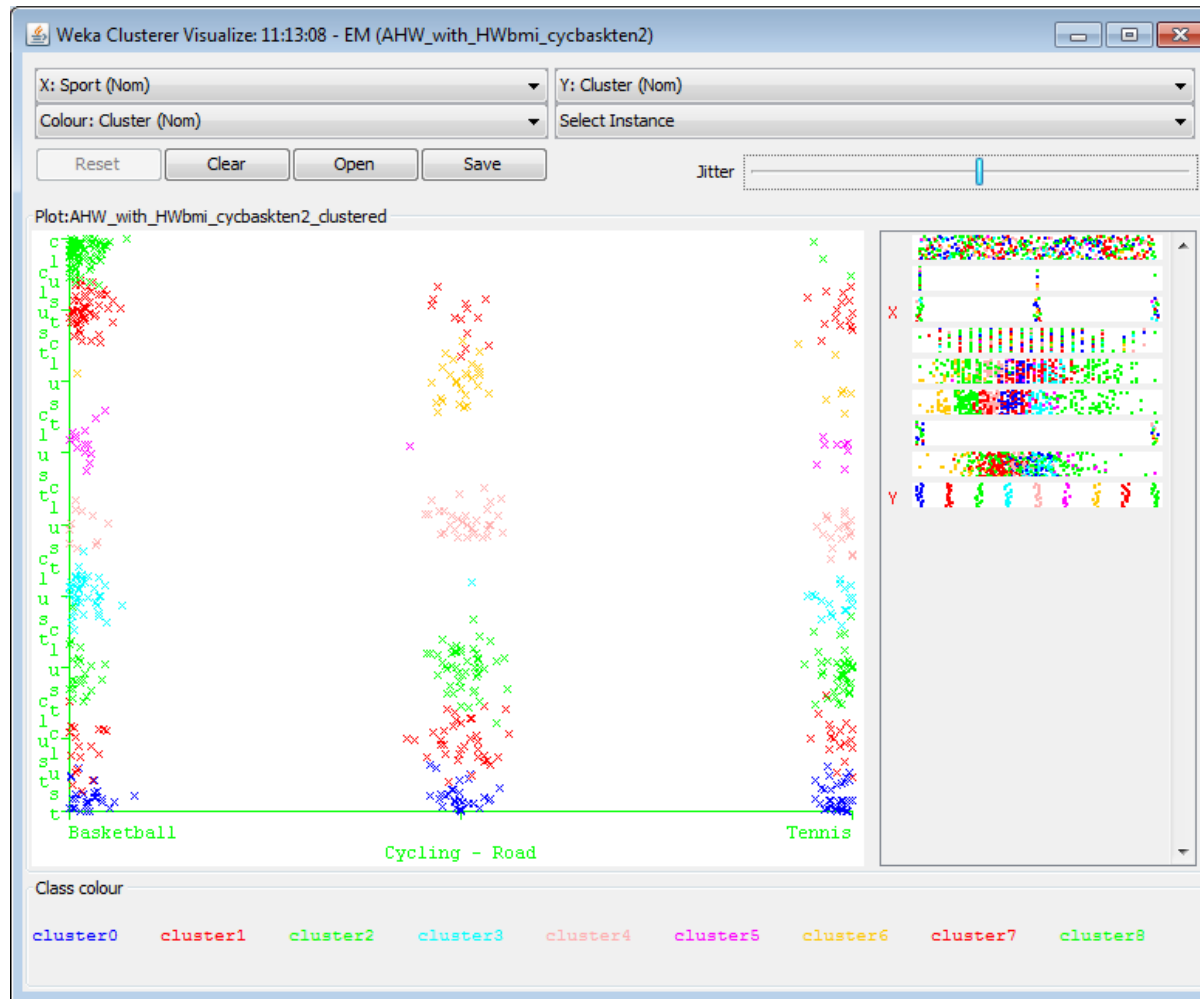
# Try EM algorithm, accept defaults and start

The screenshot shows the Weka Explorer interface with the 'Cluster' tab selected. The 'Clusterer' dropdown is set to 'EM -I 100 -N -1 -M 1.0E-6 -S 10'. The 'Cluster mode' section has 'Use training set' selected. The 'Percentage split' is set to 66%. The 'Store clusters for visualization' checkbox is checked. The 'Start' button is visible. The 'Result list' on the left shows four entries, with '11:13:08 - EM' selected. The 'Clusterer output' pane on the right displays the following table:

Attribute	Cluster 0	Cluster 1	Cluster 2	Cluster 3	Cluster 4	Cluster 5	Cluster 6	Cluster 7	Cluster 8
	(0.15)	(0.11)	(0.18)	(0.09)	(0.1)	(0.04)	(0.06)	(0.13)	(0.13)
=====									
Age									
mean	27.9888	26.9599	27.2011	27.6773	28.7174	27.0022	28.2623	26.4124	27.4742
std. dev.	3.9256	3.8879	4.6616	4.9096	4.4915	3.1316	4.0305	3.6952	3.8956
Height									
mean	183.6116	182.1607	171.1294	191.5716	176.9423	184.3566	166.9856	193.0017	205.2267
std. dev.	2.7491	2.9081	5.0946	3.0847	2.5765	6.2348	5.8676	5.2293	5.6076
Weight									
mean	77.7623	69.0197	62.3064	87.7088	70.7186	88.7604	53.6773	80.0126	105.5767
std. dev.	2.1839	2.4811	2.5788	2.8146	2.1477	5.1811	2.9723	3.5709	8.4849
HWbmi									
mean	23.0427	20.8293	21.3182	23.8995	22.6027	26.1453	19.2893	21.8768	25.0628
std. dev.	0.8296	0.6562	1.3043	0.5584	0.638	1.5337	1.3224	0.8314	1.6246
Time taken to build model (full training data) : 29.52 seconds									
=== Model and evaluation on training set ===									
Clustered Instances									
0	95 ( 15%)								

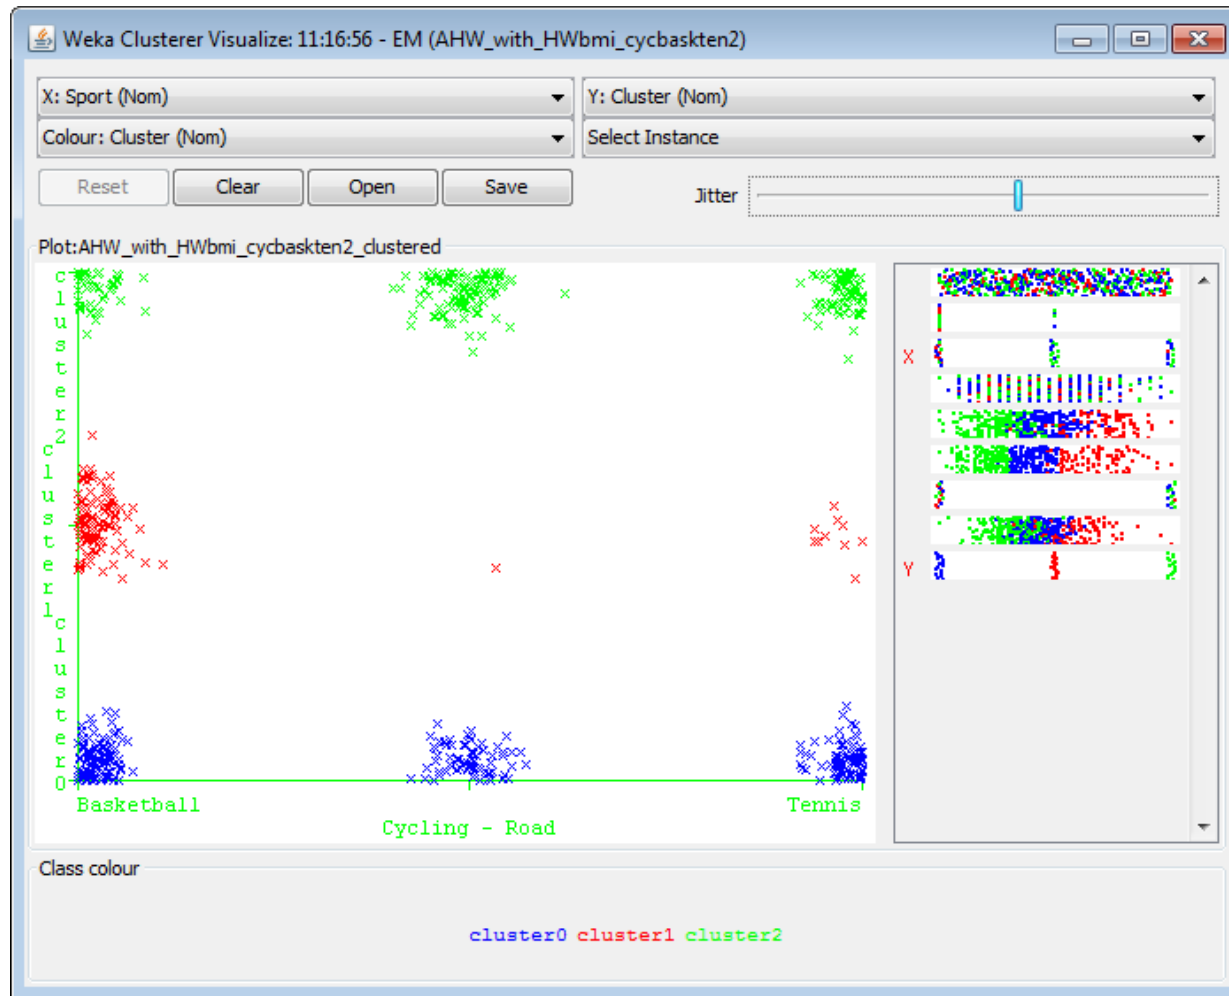
The status bar at the bottom shows 'Status OK' and a 'Log' button.

Rt click on model result, visualize cluster assignment,  
Compare sport and cluster number





# Rerun with 5 clusters, compare sport and cluster number



---

Exercise:

For simple Kmeans

Get a full sweep of K – ie change N to 1,3,4,5,  
what K would you choose?